

# vo\_\_benchmark

May 10, 2024

Notebook for two-view reconstruction without inertial data.

Author: Andre Schreiber

```
[1]: import symforce
symforce.set_epsilon_to_symbol()

import cv2
import numpy as np
import matplotlib.pyplot as plt
from pathlib import Path
from scipy.spatial.transform import Rotation as R
import time

import vo
# pose_metrics requires evo (see top of pose_metrics.py for install
  ↳ instructions)
import pose_metrics
import utils

# Note: this notebook requires pandas in addition to all of Prof. Bretl's
  ↳ dependencies
import pandas as pd
```

## 0.0.1 Read data

```
[2]: # Specify the dataset (should be 'kitti' or 'euroc')
# chosen_dataset = 'euroc'
chosen_dataset = 'kitti'

assert(chosen_dataset in ['kitti', 'euroc'])
```

## 0.0.2 Provide settings

```
[3]: # When matching (max threshold for ratio test)
if chosen_dataset == 'euroc':
    matching_threshold = 0.5
else:
    matching_threshold = 0.3

# When deciding if triangulated points are invalid
max_reprojection_err = 0.75

# Temporary folder for evo metrics
temporary_folder = Path('./temp')
temporary_folder.mkdir(parents=True, exist_ok=True)

[4]: if chosen_dataset == 'euroc': # Note: euroc takes a bit longer to load.
    # Use EuRoC MAV

    mav_video_folder = Path('./data/mav0')

    # Read MAV data
    dataset_info = utils.read_data_mav(mav_video_folder)
    print("Read dataset with keys: {}".format(sorted(list(dataset_info.
↪keys()))))

    # Extract relevant data
    cam0_K = dataset_info['cam0_K']
    cam0_distortion = dataset_info['cam0_distortion']
    visual_inertial_data = dataset_info['visual_inertial_data']

    sigma_acc_wn = dataset_info['imu_accelerometer_noise_density']
    sigma_gyr_wn = dataset_info['imu_gyroscope_noise_density']
    sigma_acc_rw = dataset_info['imu_accelerometer_noise_density']
    sigma_gyr_rw = dataset_info['imu_gyroscope_random_walk']

    dt = 1/200 # IMU frequency

    # Get extrinsics
    T_inB_ofC = dataset_info['cam0_extrinsics']
    T_inC_ofB = np.block([[T_inB_ofC[:3,:3].T, (-T_inB_ofC[:3,:3].T @
↪T_inB_ofC[:3,-1])[:3,np.newaxis]], [np.zeros(3), 1]])

    # Collate
    acc_meas, gyr_meas = utils.imu_collate(dataset_info['visual_inertial_data'])
    R_inR_ofB, v_inR_ofB, p_inR_ofB, b_a, b_w = utils.
↪groundtruth_collate(dataset_info['visual_inertial_data'], True)
```

```

    # As EuRoC's ground-truth (MoCap) is not aligned with gravity (i.e., in
    ↪ world frame), we identify the orientation of MoCap frame in world frame
    gravity = np.array([0., 0., -9.81])

    g_inB = - np.mean(acc_meas[:10], axis=0)
    g_inW = gravity

    def align_vectors(g_inB, g_inW):
        # Normalize input vectors
        g_inB_unit = g_inB / np.linalg.norm(g_inB)
        g_inW_unit = g_inW / np.linalg.norm(g_inW)

        # Compute the axis of rotation
        v = np.cross(g_inB_unit, g_inW_unit)

        # Compute the angle of rotation
        cos_theta = np.dot(g_inB_unit, g_inW_unit)
        sin_theta = np.linalg.norm(v)
        theta = np.arctan2(sin_theta, cos_theta)

        v /= np.linalg.norm(v)

        # Compute the rotation matrix
        Rot = R.from_rotvec(theta*v)
        return Rot

    # Compute rotation matrix
    R_inW_ofB = align_vectors(g_inB, g_inW)
    print("Rotation Matrix:\n", R_inW_ofB.as_matrix())

    R_inW_ofB.apply(g_inB)

    R_inW_ofR = R_inW_ofB * R_inR_ofB[:10].mean().inv()
    R_inW_ofB = R_inW_ofR * R_inR_ofB
    v_inW_ofB = R_inW_ofR.apply(v_inR_ofB)
    p_inW_ofB = R_inW_ofR.apply(p_inR_ofB)

else:
    # Use KITTI

    kitti_base_path = './data/kitti'
    kitti_date = '2011_09_26'
    kitti_drive = '0022'

    # Read KITTI data
    dataset_info = utils.read_data_kitti('./data/kitti', '2011_09_26', '0022')

```

```

# Extract relevant data
cam0_K = dataset_info['cam0_K']
cam0_distortion = dataset_info['cam0_distortion']
visual_inertial_data = dataset_info['visual_inertial_data']

T_inC_ofB = dataset_info['cam0_extrinsics']
R_inB_of_C = T_inC_ofB[:3, :3].T
t_inB_of_C = R_inB_of_C @ T_inC_ofB[:3, 3]
T_inB_ofC = np.block([[R_inB_of_C, t_inB_of_C[:3, np.newaxis]], [np.zeros(3), 1]])

R_inC_of_B = T_inC_ofB[:3, :3]
t_inC_of_B = T_inC_ofB[:3, 3]

# Collate
acc_meas, gyr_meas = utils.imu_collate(visual_inertial_data)
R_inW_ofB, v_inW_ofB, p_inW_ofB, b_a, b_w = utils.
groundtruth_collate(visual_inertial_data, False)

```

### 0.0.3 Create random generator

```
[5]: rng = utils.create_rng(42)
```

seeding RNG with 42

### 0.0.4 Create image keypoint feature extractor

```
[6]: feature_extractor = cv2.SIFT_create() # could also do ORB_create() for ORB
features
```

### 0.0.5 Two view reconstruction

Get initial solution

```
[7]: if chosen_dataset == 'euroc': # Note: euroc takes a bit longer to load.
    # Use EuRoC MAV
    chosen_index = 500
    advance = 100
else:
    chosen_index = 50
    advance = 5

# Get first index closest to chosen index

```

```

first_frame_idx = utils.get_index_of_next_image(visual_inertial_data,
    ↪chosen_index)
# Get second index
second_frame_idx = utils.get_index_of_next_image(visual_inertial_data,
    ↪first_frame_idx+advance)

# Create two views
views = [
    vo.create_view_data(utils.
    ↪read_image(visual_inertial_data[first_frame_idx]['image_file']),
        first_frame_idx, feature_extractor, cam0_K,
    ↪cam0_distortion),
    vo.create_view_data(utils.
    ↪read_image(visual_inertial_data[second_frame_idx]['image_file']),
        second_frame_idx, feature_extractor, cam0_K,
    ↪cam0_distortion)
]

# Perform two-view reconstruction
tic = time.time()
tracks = vo.vo_2view(views, matching_threshold, cam0_K, rng, use_opencv=False)
toc = time.time()

analytical_guess = toc - tic
print(f"Analytical guess: {analytical_guess:.2f} [s]")

```

found 145 good matches  
 found 145 inliers  
 Analytical guess: 2.84 [s]

```

[8]: vo.show_reproj_results(views, tracks, cam0_K, cam0_distortion,
    ↪print_raw_reproj=True, show_reproj_histogram=True)
vo.visualize_predictions(views, tracks, cam0_K, cam0_distortion)

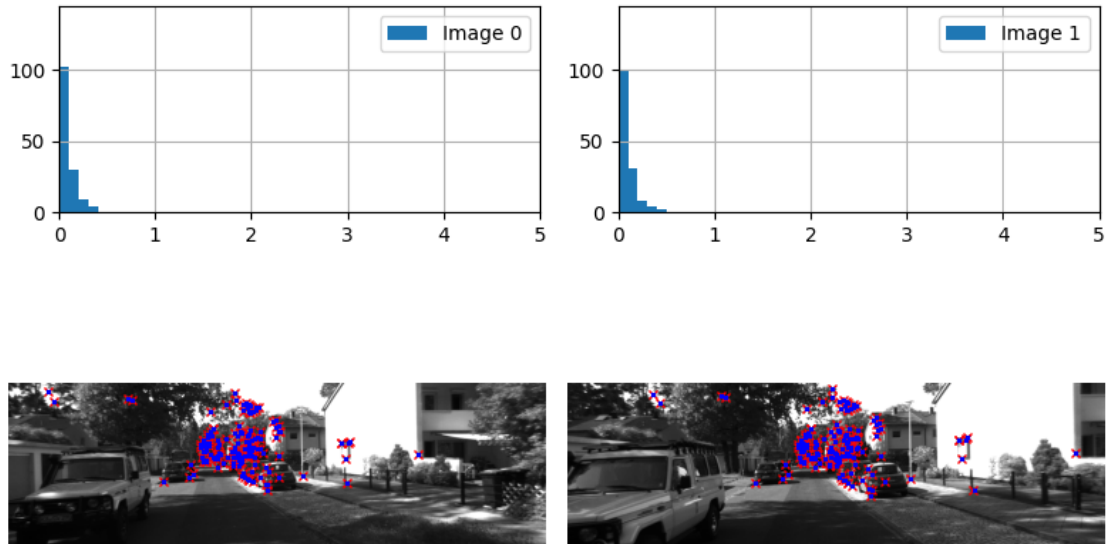
```

#### REPROJECTION ERRORS

```

Image 0 ( 145 points) : (mean, std, max, min) = (0.0868, 0.0800, 0.4026,
0.0004)
Image (raw reprojection) 0 ( 145 points) : (mean, std, max, min) = (0.0868,
0.0800, 0.4026, 0.0004)
Image 1 ( 145 points) : (mean, std, max, min) = (0.0931, 0.0883, 0.4769,
0.0004)
Image (raw reprojection) 1 ( 145 points) : (mean, std, max, min) = (0.0931,
0.0883, 0.4769, 0.0004)

```



Get post-optimization solution

Run below to keep the initial views and tracks

```
[9]: views_ini = views.copy()
     tracks_ini = tracks.copy()
```

```
[10]: tic = time.time()

views, tracks, initial_values, results = vo.vo_nonlinear_optimize(views_ini,
↳ tracks_ini, cam0_K, max_reprojection_err)

toc = time.time()
nonlinear = toc - tic
print(f"{nonlinear:.2f} [s]")
```

```
[2024-05-10 16:46:52.493] [info] LM<sym::Optimize> [iter    0] lambda:
1.000e+00, error prev/linear/new: 6.158e+00/0.000e+00/6.046e+00, rel reduction:
1.82341e-02
[2024-05-10 16:46:52.519] [info] LM<sym::Optimize> [iter    1] lambda:
1.000e-01, error prev/linear/new: 6.046e+00/0.000e+00/5.926e+00, rel reduction:
1.99465e-02
[2024-05-10 16:46:52.539] [info] LM<sym::Optimize> [iter    2] lambda:
1.000e-02, error prev/linear/new: 5.926e+00/0.000e+00/5.894e+00, rel reduction:
5.30232e-03
[2024-05-10 16:46:52.558] [info] LM<sym::Optimize> [iter    3] lambda:
1.000e-03, error prev/linear/new: 5.894e+00/0.000e+00/5.858e+00, rel reduction:
6.04853e-03
[2024-05-10 16:46:52.578] [info] LM<sym::Optimize> [iter    4] lambda:
```

```

1.000e-04, error prev/linear/new: 5.858e+00/0.000e+00/5.760e+00, rel reduction:
1.68508e-02
[2024-05-10 16:46:52.598] [info] LM<sym::Optimize> [iter    5] lambda:
1.000e-05, error prev/linear/new: 5.760e+00/0.000e+00/5.622e+00, rel reduction:
2.39643e-02
[2024-05-10 16:46:52.618] [info] LM<sym::Optimize> [iter    6] lambda:
1.000e-06, error prev/linear/new: 5.622e+00/0.000e+00/5.561e+00, rel reduction:
1.07297e-02
[2024-05-10 16:46:52.637] [info] LM<sym::Optimize> [iter    7] lambda:
1.000e-07, error prev/linear/new: 5.561e+00/0.000e+00/5.553e+00, rel reduction:
1.47557e-03
[2024-05-10 16:46:52.657] [info] LM<sym::Optimize> [iter    8] lambda:
1.000e-08, error prev/linear/new: 5.553e+00/0.000e+00/5.552e+00, rel reduction:
1.28544e-04
[2024-05-10 16:46:52.677] [info] LM<sym::Optimize> [iter    9] lambda:
1.000e-09, error prev/linear/new: 5.552e+00/0.000e+00/5.552e+00, rel reduction:
6.95137e-06
0.84 [s]

```

```

[11]: vo.show_reproj_results(views, tracks, cam0_K, cam0_distortion,
    ↪ print_raw_reproj=True, show_reproj_histogram=True)
vo.visualize_predictions(views, tracks, cam0_K, cam0_distortion)


```

#### REPROJECTION ERRORS

```

Image 0 ( 144 points) : (mean, std, max, min) = (0.0745, 0.0797, 0.5337,
0.0001)
Image (raw reprojection) 0 ( 144 points) : (mean, std, max, min) = (0.0745,
0.0797, 0.5337, 0.0001)
Image 1 ( 144 points) : (mean, std, max, min) = (0.0699, 0.0719, 0.4287,
0.0001)
Image (raw reprojection) 1 ( 144 points) : (mean, std, max, min) = (0.0699,
0.0719, 0.4287, 0.0001)

```





```
[12]: print(f"\nTiming\n{'='*100}")
      print(f"Analytical guess: {analytical_guess:.2f} [s]")
      print(f"Non linear (V0): {nonlinear:.2f} [s]")
```

Timing

```
=====
```

```
=====
Analytical guess: 2.84 [s]
Non linear (V0): 0.84 [s]
```

```
[13]: # This highlights the scale ambiguity seen in two-view reconstruction
      print("Metric pose difference norm (gt) = {:.3f}".format(np.linalg.norm(
          # Note: while this is in world frame, the frames do not have any scaling,
          # so measuring displacement this way still should provide the correct
          ↪distance.
          np.linalg.norm(p_inW_ofB[first_frame_idx] - p_inW_ofB[second_frame_idx])
      )))
      print("Metric pose difference norm (ini) = {:.3f}".format(np.linalg.
          ↪norm(results.initial_values['T_inB1_ofA'].t)))
      print("Metric pose difference norm (sf) = {:.3f}".format(np.linalg.norm(results.
          ↪optimized_values['T_inB1_ofA'].t)))
```

```
Metric pose difference norm (gt) = 2.784
Metric pose difference norm (ini) = 1.000
Metric pose difference norm (sf) = 1.000
```

Relative change in position ( $\Delta p$ ) and rotation ( $\Delta R$ ) between frames (**NOT ERROR w.r.t Ground Truth**)

```
[14]: # First frame results
      frame0_key = 'T_inB0_ofA'
      # Symforce - Initial values
      R_inB0_ofA_ini = results.initial_values[frame0_key].R.to_rotation_matrix()
      p_inB0_ofA_ini = results.initial_values[frame0_key].t

      # Symforce - Optimized values
      R_inB0_ofA_sf = results.optimized_values[frame0_key].R.to_rotation_matrix()
      p_inB0_ofA_sf = results.optimized_values[frame0_key].t
```



```

# ground truth
R_inW_ofB0_gt = R_inW_ofB[first_frame_idx].as_matrix()
p_inW_ofB0_gt = p_inW_ofB[first_frame_idx]
v_inW_ofB0_gt = v_inW_ofB[first_frame_idx]

# Second frame results
frame1_key = 'T_inB1_ofA'
# Symforce - Initial values
R_inB1_ofA_ini = results.initial_values[frame1_key].R.to_rotation_matrix()
p_inB1_ofA_ini = results.initial_values[frame1_key].t

# Symforce - Optimized values
R_inB1_ofA_sf = results.optimized_values[frame1_key].R.to_rotation_matrix()
p_inB1_ofA_sf = results.optimized_values[frame1_key].t

# Ground Truth
R_inW_ofB1_gt = R_inW_ofB[second_frame_idx].as_matrix()
p_inW_ofB1_gt = p_inW_ofB[second_frame_idx]
v_inW_ofB1_gt = v_inW_ofB[second_frame_idx]

print(f"RELATIVE CHANGE BETWEEN FRAMES (NOT ERROR)\n")
print(f'dp\n{"="*50}')
print('(Analytical guess) dp: {:.2f} [m]'.format( np.linalg.norm(
    ↪ p_inB0_ofA_ini - p_inB1_ofA_ini)))
print('(Non linear - VO) dp: {:.2f} [m]'.format( np.linalg.norm( p_inB0_ofA_sf ↪
    ↪ - p_inB0_ofA_sf)))
print('(Ground Truth)      dp: {:.2f} [m]'.format( np.linalg.norm( p_inW_ofB0_gt↪
    ↪ - p_inW_ofB1_gt)))

print(f'\ndR\n{"="*50}')
print('(Analytical guess) dR scalar: {:.5f} [deg]'.format(pose_metrics.
    ↪ rotational_error( R_inB0_ofA_ini, R_inB1_ofA_ini)))
print('(Non linear - VO) dR scalar: {:.5f} [deg]'.format(pose_metrics.
    ↪ rotational_error( R_inB0_ofA_sf, R_inB1_ofA_sf)))
print('(Ground Truth)      dR scalar: {:.5f} [deg]'.format(pose_metrics.
    ↪ rotational_error( R_inW_ofB0_gt, R_inW_ofB1_gt)))

```

RELATIVE CHANGE BETWEEN FRAMES (NOT ERROR)

dp

```

=====
(Analytical guess) dp: 1.00 [m]
(Non linear - VO) dp: 0.00 [m]
(Ground Truth)      dp: 2.78 [m]

```

dR

=====

(Analytical guess) dR scalar: 7.10949 [deg]

(Non linear - V0) dR scalar: 7.10223 [deg]

(Ground Truth) dR scalar: 7.21199 [deg]